

PulseDNA: AI Trust Infrastructure — Presenter Script

Target Audience: Andrew Ng & Technical ML Investors

Pacing: ~20 seconds per slide

Psychological Goal: Guide the listener from *"AI has a credibility crisis"* → *"PulseDNA is the SSL of AI"* → *"This is a \$10M infrastructure build"* (without ever saying the number). Secure intellectual respect first; capital will follow.

Slide 01: Building the Trust Infrastructure for the AI Era

"Andrew, thank you for the time. We believe we are at a structural inflection point. AI answers are rapidly replacing traditional search. But the infrastructure to actually *verify* those answers doesn't exist yet. We are building that missing layer — the trust infrastructure for the AI era."

Slide 02: AI Has a Trust Problem

"Right now, the world's most powerful models — from OpenAI to Gemini to LLaMA — all share the same fatal flaw: they generate highly confident answers with zero verifiable provenance. In high-risk sectors like nuclear, finance, or intelligence, a hallucination isn't just an annoyance; it's a catastrophic liability."

Slide 03: The Hidden Cost of Unverified AI

"The market is already feeling this pain financially. When AI can't be trusted, enterprises are forced to put expensive humans in the loop just to double-check the machine. The ROI of AI adoption is being eaten alive by the manual cost of verification. Trust isn't just a safety issue; it's an economic bottleneck."

Slide 04: Why Current AI Is Structurally Broken

"You know better than anyone that LLMs are probabilistic token predictors, not factual reasoning engines. They cannot self-correct their own provenance. You can't prompt-engineer your way out of this. The solution isn't a better model; the solution is an external, deterministic infrastructure layer."

Slide 05: Magic Lens Trust Infrastructure

"This is where PulseDNA comes in. We're building a layered, event-driven architecture. To proactively answer a question we get a lot — yes, this is a streaming architecture, not batch

processed. We use Kafka-style pipelines to handle millions of signals per day in real-time, operating at the same foundational layer as a Databricks or a Confluent."

Slide 06: The Provenance Engine

"At the core is our Provenance Engine. If you're wondering how we compute credibility scores: it's a deterministic mix of graph relationships, citation chains, and publisher credibility signals. We map the lineage of every claim back to its root source, creating a verifiable mathematical proof of trust."

Slide 07: Why This Moat Is Defensible

"This creates a massive data flywheel. As we ingest more sources, the trust graph gets richer. A richer graph means more accurate credibility scores, which attracts more enterprise customers, who in turn feed more proprietary signals back into the network. The graph itself becomes an insurmountable moat."

Slide 08: A Global Trust Signal Network

"Zooming out, what we are really building is a global credibility graph for the internet. The internet solved identity with SSL in 1994. We are building the equivalent for AI credibility in 2026. *I'd really value your perspective here — if you were building trust infrastructure for AI today, what would you prioritize first?*"

Slide 09: Traction & Proof Points

"We aren't just theorizing; we are deploying. We're recognized by the Google AI Futures Fund, partnered with NVIDIA, and actively deploying with Hadron Energy and Supermicro. We are proving this architecture works in the most unforgiving regulatory environment on earth."

Slide 10: Start With High-Trust Industries

"Our GTM motion starts where verified knowledge is non-negotiable: Nuclear Energy. From there, we expand into financial analysis, research, and government intelligence. Compliance is our beachhead. Once we are embedded in an enterprise's compliance workflow, we cease being software and become critical infrastructure."

Slide 11: Why Now

"The timing is forced by regulation. The EU AI Act, the NRC, and the SEC are no longer asking for AI transparency; they are mandating it. The window to establish the standard for AI provenance is open right now, and the market is desperate for a deterministic solution."

Slide 12: Strategic Partners

"We are going to market alongside the hardware and data leaders of the space. Supermicro for sovereign, air-gapped compute. NVIDIA for acceleration. John Snow Labs for clinical and unstructured data integration. We are slotting directly into the existing enterprise AI stack."

Slide 13: The Minds Behind Magic Lens

"Finally, the team. I honestly feel the universe has been steering us here. When we landed in DOE R&D, we were gifted this 'dream team' of experts in NLP, physics, and nuclear sensor data. Between myself, Eugene, and Loic, we have the exact cross-section of Deep Tech, AI, and nuclear engineering required to build this. We'd love to have you along for the journey."

Appendix: Q&A Preparation for Andrew Ng

The goal here is not to get an immediate 'yes', but to secure his intellectual respect. If he walks away thinking "this is a real infrastructure idea," he will follow your progress, introduce you to other investors, and potentially invest later.

1. "How do you compute trust scores?"

Your Answer: "It's a deterministic combination of graph relationships, citation chains, and publisher credibility signals. We don't rely on LLM vibes; we rely on verifiable mathematical lineage back to the root source."

2. "What data do you start with?"

Your Answer: "We start with the public web, academic research repositories, government datasets like the NRC and SEC, and direct publisher feeds. The nuclear beachhead gives us a highly structured, bounded dataset to perfect the engine."

3. "What's the smallest working loop?"

Your Answer: "Source ingestion → Verification → Trust score generation → API response. That thin slice is what we demo to enterprise clients to prove the architecture works before we scale it."

4. "Why hasn't Google done this?"

Your Answer: "Because their objective function is different. Search engines optimize for relevance, ranking, and engagement. We are optimizing purely for verifiable provenance and credibility. It's a fundamentally different engineering problem." *(Note: This answer is incredibly strong with ML thinkers).*

5. "How large could the trust graph become?"

Your Answer: "At internet scale, it becomes the definitive credibility layer for all digital information. Just as Google indexed the web, PulseDNA verifies it. The graph scales horizontally to encompass billions of nodes."